



Evaluating Resiliency and Performance of Networked Satellite Systems

Evaluierung der Resilienz und Leistung von Vernetzten
Satelliten Systemen

Robert Alexander Uwe Richter

Universitätsmasterarbeit
zur Erlangung des akademischen Grades

Master of Science
(M. Sc.)

im Studiengang
IT Systems Engineering

eingereicht am 11. November 2024 am
Fachgebiet Data-Intensive Internet Computing der
Digital-Engineering-Fakultät
der Universität Potsdam

Gutachter

Prof. Dr. Vaibhav Bajpai
Prof. Dr. Tilmann Rabl

Betreuer

Dr. Vasilis Ververis

Abstract

Acknowledgments

Contents

Abstract	iii
Acknowledgments	v
Contents	vii
1 Introduction	1
1.1 Motivation	1
1.2 Research Questions	1
2 Background	3
2.1 Satellite Communication Explained	3
2.1.1 Bent-Pipes	4
2.2 Usecase of Networked Satellite Systems	4
2.3 Networked Satellite ISPs	5
2.4 Satellite Network Simulators	6
2.4.1 StarPerf	6
2.4.2 Hypatia	6
2.4.3 Problems of Network Simulators	7
2.5 Growth of Satellite Constellations	7
3 Methodology	9
3.1 Data Collection	9
3.2 Reproducibility of Results and Data	11
3.3 Measurement Platforms	11
3.3.1 RIPE Atlas	11
3.3.2 Cloudflare Radar	12
3.3.3 OONI	13
3.3.4 N2YO	13
4 Results	15
4.1 Latency over whole Data Range	15
4.2 Latency Development on Individual Days	18

4.3	Latency Development on Individual Months	18
4.3.1	Latencies on Different Weekdays	18
4.4	Disconnect Events of RIPE Atlas Probes	20
4.4.1	Correlation of Disconnect Events with other Metrics . . .	23
4.5	Latency and Packet Loss Correlation	23
4.5.1	Correlation in Single Years	23
4.5.2	Correlation with the Number of Probes	26
4.6	Traceroute Analysis	27
4.6.1	Privacy Concerns in Traceroute Data	28
5	Conclusions & Outlook	31
6	Appendix	33
	Bibliography	37
	Declaration of Authorship	39

1.1 Motivation

1.2 Research Questions

In the following thesis, we want to solve the following research questions.

Link mentioned platforms to chapters of the individual measurement platforms.

Research Question 1: How do networked satellite systems perform in terms of latency and packet loss?

Packet loss and latency are the main factors for the performance of a network. Recent research stated high packet loss values, but mostly idealized latency values. This thesis shall take a look at those two performance characteristics. To solve this research question, we will use RIPE Atlas and Cloudflare Radar.

Research Question 2: Do latency and packet loss correlate?

It is expected that a rising latency correlates with a rising packet loss and vice versa due to the need of retransmits. It is in question whether this is the case for networked satellite systems.

If it appears that there is no clear correlation observable, what factors might influence the correlation?

To solve this research question, we will use the data obtained from RQ 1 (i.e., data from RIPE Atlas and Cloudflare Radar).

Research Question 3: How do networked satellite systems route?

Networked satellite systems face a severe challenge when routing a packet. Additionally, to the routing problem, they need to find an ideal route through their satellite constellation involving sending data to a ground station.

This thesis takes a look at the routing behavior of networked satellite systems. To solve this research question, we will use traceroute data from RIPE Atlas.

2.1 Satellite Communication Explained

A satellite communication system has the target to send packets from a sender connected via a networked satellite ISP to a receiver connected to a terrestrial ISP.




Figure 2.1: Schematics of Satellite Communication [Moh+24].

The structure is illustrated in Figure 2.1. Initially, there is a sender connected to a satellite antenna provided by the specific networked satellite ISP. The satellite antenna is also colloquially called "dishy". From the senders' perspective, there is no different use compared to a traditional terrestrial internet connection. The antenna will send packets to a satellite that might either forward the packets to another satellite or directly send them to a nearby Ground Station (GS) (also called Point of Presence (PoP)). The GS serves as an "entry point" to the terrestrial internet. A GS can forward packets through wired connections, so they will reach the receiver eventually.

The technology of communicating with satellites will not be explained here. Amongst others, it is described by Pratt et al. [PA19].

Theoretically, the propagation of data in satellite communication happens at the speed of light. However, the latencies still vary a lot depending on the choice of networked satellite ISP. Each of them provides a different satellite constellation. One of the dominating factors is the altitude of the constellation. The higher the satellites are positioned, the more region a single satellite covers. Therefore, fewer satellites are required to cover a specific region, which reduces costs.

Overall, one differentiates between Geostationary Orbit (**GEO**) (35'786 km altitude), Medium-Earth Orbit (**MEO**) (2'000 - 35'786 km altitude), and Low-Earth Orbit (**LEO**) (< 2'000 km altitude) satellite constellations. One can see that this has a significant impact on the latency.

$$\frac{2 \cdot 35'786 \text{ km}}{300'000 \frac{\text{km}}{\text{s}}} \approx 0.240 \text{ s} \quad (2.1)$$

$$\frac{2 \cdot 550 \text{ km}}{300'000 \frac{\text{km}}{\text{s}}} \approx 0.004 \text{ s} \quad (2.2)$$

A **GEO** constellation has a minimum latency of around 240 ms, as shown in Equation 2.1, assuming a packet needs to reach a satellite in GEO altitude and get back to the earth's surface. On the other hand, Equation 2.2 shows that a **LEO** satellite constellation at an altitude of 550 km (like in the case of Starlink) has a minimum latency of only 4 ms. Research has shown that even in practice **LEO** constellations are superior compared to **GEO** constellations, especially in terms of latency [Ram+23; Seg20]. However, terrestrial internet performed better compared to **LEO** constellations [Ma+23; Moh+24].

2.1.1 Bent-Pipes

Ground stations are required for satellites to route packets to a receiver. Depending on the location of the sender, the first satellite on the route might have to forward packets to a different satellite (using an Inter-Satellite Link (**ISL**)) in order to reach a region with a **GS**. Such a route is called "Bent-Pipe". Depending on the number of satellites on the route, it is called an "n-hop-bent-pipe". If there is only a single satellite involved, it is a 1-hop-bent-pipe.

It is in question how the bent-pipe influences the performance and behavior of a networked satellite system. Research still discusses whether bent-pipes provide a positive ([Hauri2020]) or negative ([Moh+24]) impact.

2.2 Usecase of Networked Satellite Systems

Networked satellite systems are around for quite a while since humanity possesses both satellites and the internet. The internet has become a key-technology for communication in nearly every aspect of the society. Having no access to the internet will result in significant drawbacks. However, internet requires a complex

Table 2.1: Different networked satellite ISPs

ISP	Category	Customer Group
Starlink	LEO	Private & Business
OneWeb	LEO	Business
HughesNet	GEO	Private & Business
Intelsat	GEO	Business
Viasat	GEO	Business
Orbcomm	GEO	Business
Iridium	GEO	Private & Business

infrastructure with high cost. Especially in distant location with little population, building such an infrastructure will not be affordable.

Therefore, people came up with the idea of using satellites to communicate with distant locations. In theory, only three satellites are required to communicate with any point on earth. The cost of providing this number of satellites is much less compared to the costs of providing a terrestrial network infrastructure with similar accessibility.

There is one group of people that will not receive terrestrial network infrastructure: people on boats and planes. While there is cellular internet access, it is not offered on the sea and in the high sky. Therefore, the only option of internet access is satellite internet, which is served all over the world.

Additionally, networked satellite systems are much more resilient to physical influences like earthquakes, terrorist attacks, or storms [Ste+24]. Terrestrial infrastructure can be destroyed easily and therefore especially governments contracted networked satellite system ISPs.

2.3 Networked Satellite ISPs

There are a couple of ISPs out there providing internet access via satellites.

Table 2.1 shows different ISPs for networked satellite systems. The differentiate mostly by providing a LEO or GEO service. The only LEO ISPs are Starlink and OneWeb. Starlink has an economical advantage over most of its competitors as they are attached to SpaceX. SpaceX is a company providing spacecraft manufacturing and launch services. This makes the launch of satellites less tedious compared to Starlink's competitors.

Kohnmann [Koh24] reported about intentions of Amazon launching a networked satellite provider called Kuiper. It is planned to be launched by the end of 2026.a

2.4 Satellite Network Simulators

Amongst concrete measurements, one can also simulate networked satellite systems. This became increasingly interesting when the constellations were composed of many more satellites compared to traditional GEO satellite constellations. For example, LEO constellations comprise hundreds to thousands of satellites, which implies a highly complex system.

Sadly, measurements are often highly difficult as they either require acquiring satellite hardware or recruiting users that already possess the required hardware. Simulation would tackle both problems, while maintaining low cost. To the best of our knowledge, we found two networked satellite simulators for LEO constellations.

2.4.1 StarPerf

*StarPerf*¹ [LLL20] is a mega-constellation performance simulation platform. It specifically aims at measuring the impact of the movements of satellites. Also, it measures performance in different areas. However, setting it up required, amongst others, Matlab and STK. This made the project difficult and expensive to test. Therefore, we did not advance in trying out *StarPerf*.

2.4.2 Hypatia

*Hypatia*² [Kas+20] is another LEO network simulation framework, released in 2020 just like *StarPerf*. It aims at a low-level simulation on packet-level and visualizes the data. Unlike *StarPerf*, it only requires a Python3 installation. Sadly, running simulations with Hypatia is highly complex as it requires the user to define, amongst others, the satellites, ground stations, and points of presence. This information is hardly available, which renders the simulations barely usable. Also Hypatia is not maintained anymore.

1 [SpaceNetLab/StarPerf_Simulator](#)

2 [snkas/hypatia](#)

2.4.3 Problems of Network Simulators

To the best of our knowledge, research has stopped relying on simulators since 2020. There is proper hardware available that allows testing in the real world. Testing in the real world has the advantage as it takes more variables into account. Crucial factors for the performance of a networked satellite system are the weather, congestion, solar magnetic storms, material failure, and many more. Those cannot be ideally tested with simulations and will eventually produce wrong results. Therefore, for further research of this thesis, simulations will not be used.

2.5 Growth of Satellite Constellations

Recent years saw a rapid development of satellite technology, especially due to the growing demand of global connectivity and communication. Therefore, companies constructed their own satellite constellations leading to a total number of more than 29'000 objects in space at the time of writing.




Figure 2.2: Growth of number of satellites in different satellite constellation from 2000 to 2024 (note the logarithmic scale).

Figure 2.2 shows various satellite constellations with the number of satellites per year. Table 2.2 show the corresponding numbers, starting in 2017. One can see that Starlink is by far the constellation with most satellites. At the time of writing, it has 6'396 satellites. Starlink grew from 2023 to 2023 by nearly 2000 satellites. OneWeb,

Table 2.2: Growth of satellite constellations from 2017 to 2024

	2017	2018	2019	2020	2021	2022	2023	2024
Classification								
Starlink	0	0	120	943	1871	3481	5326	6396
Orbcomm	63	63	63	63	63	63	63	61
OneWeb	2	2	3	104	388	498	628	628
Beidou	15	31	38	39	39	39	41	41
Westford Needles	47	47	46	46	45	45	44	44
O3B Networks	12	16	20	20	20	20	20	20
Intelsat	74	74	74	74	74	74	74	74

Starlink’s closest competitor, has a total of 628 satellites with no change between 2023 and 2024.

3.1 Data Collection

For analysis, we create a dataset containing measurements from various sources about networked satellite systems. Sadly, only data from Starlink devices was included in the thesis as other satellite network systems did not have sufficient data openly available.

The data originates from RIPE Atlas, Cloudflare Radar, OONI, and N2YO. While N2YO provides data about satellites, the others hold measurement results. In the case of N2YO, the website was crawled, while the others provide API access.

The process of data collection is illustrated in Figure 3.1.

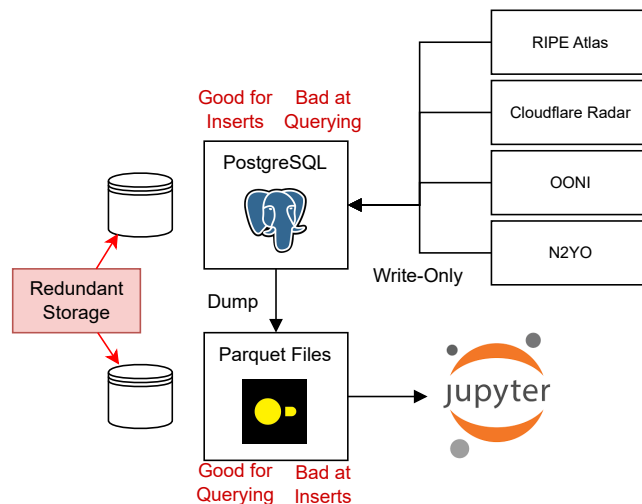


Figure 3.1: Architecture of Data Collection

The data is collected from each platform and inserted into a PostgreSQL database. The database shall allow producers to quickly insert new data (i.e., new rows). Transactional databases are the best choice for that, e.g., PostgreSQL. To quickly analyze data, an analytical database is the best choice. For that purpose, Parquet

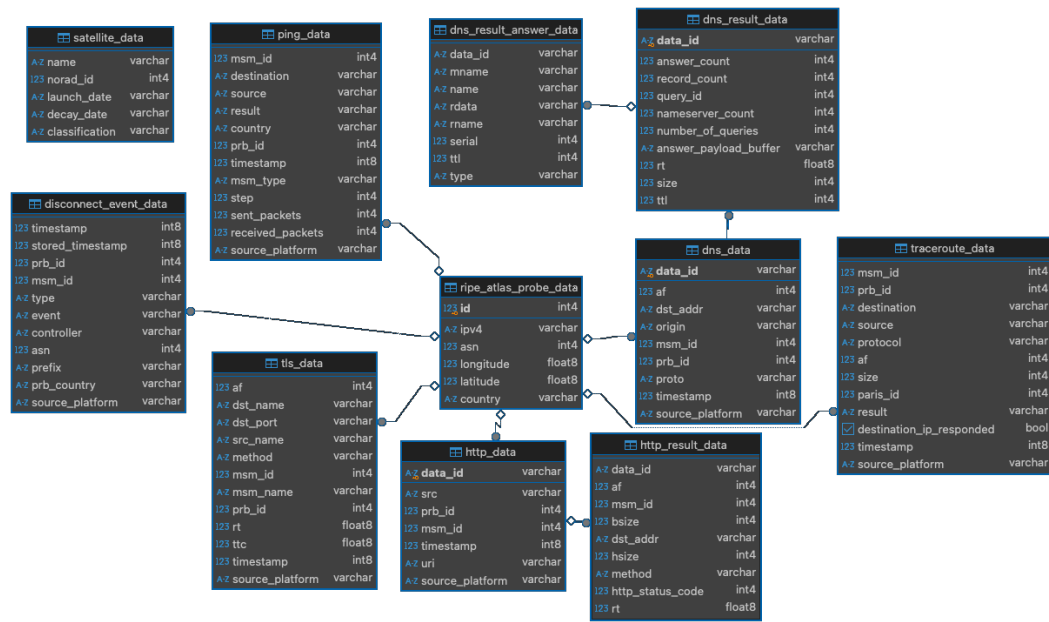


Figure 3.2: ER Diagram on Data Schema in PostgreSQL Database.

files can be used. Therefore, the data from PostgreSQL is dumped into the Parquet files. This creates a redundant storage, for the sake of analysis speed.

The resulting data format is shown in Figure 3.2.

The database includes Ping, Traceroute, TLS, HTTP, Disconnect Event, and DNS measurement data, as well as information about the RIPE Atlas probes and all satellites ever launched (including rocket bodies). The whole database comprises more than 150 GB.

Additionally, the analysis uses data from IPinfo. This data is however not included in the dataset and has to be obtained from IPinfo itself.

RIPE Atlas Data

RIPE Atlas offers various probes connected via networked satellite systems. At the moment of writing, all of those probes are connected via Starlink. Therefore, all the data from RIPE Atlas probes is Starlink data.

Probes are computers running the [probe software from RIPE Atlas](#). They are centrally connected to the RIPE Atlas servers and can be used by any person to perform measurements against them. The possible measurements are defined by the [RIPE Atlas REST API](#).

Overall, there are 150 probes from 26 countries. Each probe performs basic measurements on a regular schedule (so-called built-in measurements). The built-in measurements are the main source of data and serve as historical record. This allows to analyze data from 2022. Even if there is Starlink data prior to 2022, it originates from very little probes and therefore will not be considered in this thesis to avoid unreliable data.

3.2 Reproducibility of Results and Data

The code to reproduce the results and data is uploaded to GitHub³. The user is only required to have a running internet connection and an installation of Nix and Docker. Follow the instructions provided by the repository to obtain the results and data by yourself⁴.

3.3 Measurement Platforms

The following chapter describes the measurement platforms that were used to measure and collect data.

3.3.1 RIPE Atlas

RIPE Atlas is an open measurement platform to perform basic measurements on the internet. It holds more than ten thousand probes that serve as start points. A probe is machine that has the RIPE Atlas probe software installed. A user can request a probe for a measurement (e.g., a ping from a probe to a RIPE NCC root server).

RIPE Atlas offers the following basic types of measurements:

- Ping
- Traceroute
- (Dis)connection Events
- DNS Lookup
- DNS TLDs

³ See <https://github.com/starlink-thesis-diic/starlink-thesis>

⁴ The scripts might not work, if data format in the sources has changed since the last update of the scripts.

- HTTP
- TLS (SSL) GET Certificate

RIPE Atlas offers the possibility to start measurements via the web interface or an API. While the web interface is sufficient for most use cases, it is also quite limited. Using the API offers more possibilities (e.g., turning DoH queries into DoTLS by adding `"tls": true` into the definition of the measurement).

Additionally, each registered probe performs measurements on a regular basis in fixed time intervals (e.g., each probe performs every 240 seconds a ping measurement against all RIPE NCC root servers). Those measurements are called built-in measurements. They run always when the probe is only and connected to the RIPE Atlas network. The built-in measurements are the ones used for the data of this thesis within a specified time interval (usually January 2022 to June 2024).


Downloading the results of a measurement also works by accessing the API. For previous results, one can access the REST API that stores all results ever measured. For more up-to-date results, one can use the Streaming API that sends the most recent results to a subscribed socket.

Theoretically, the API is rate-limited, but we did not encounter problems even when bulk-downloading larger datasets. Still, ethical crawling was used.

3.3.2 Cloudflare Radar

Cloudflare manages a vast amount of traffic on the internet. Initially, it served the purpose of protecting servers from attacks. Today amongst others, it also manages traffic, and sells data. In 2020, Cloudflare released data that has been held internal. This new platform is Cloudflare Radar ([CFR](#)). [CFR](#) offers various statistics on the internet ranging from security statistics over latency measurements to bot ratios per country.

For the purpose of this thesis, we are mostly interested in the performance measurements. [CFR](#) offers performance measurements within their Internet Quality Index ([IQI](#)). It is the collection of measurements performed through their [Internet Speed Test](#) [[Clo](#); [Dav23](#)]. As target, [CFR](#) uses a fixed set of Cloudflare servers.



(a) Locations of Root Servers (all *.root-servers.org together) [roo24] (b) Locations of Cloudflare Data Centers [Clo24]

Takeaway: Difference of RIPE Atlas and Cloudflare Radar

Measurements performed in CFR and RIPE Atlas differentiate in the targeted servers (root servers vs. Cloudflare servers), the measurement probe (RIPE Atlas probe vs. browser), and the metric used (TLS vs. Time to First Byte). Cloudflare's servers are expected to have a better latency (they claim to be 50ms away from 95 % of the earth's population; State September 17, 2024).

3.3.3 OONI

3.3.4 N2YO

N2YO is a platform collecting data about satellites in space. It is a data integration platform collecting data from different platforms like Celestrak and Space-Track.

For the purpose of this thesis, N2YO is used to find information on the development of various satellite constellations. Therefore, we want to find the following information:

- Satellite ID (also formerly called Norad ID)
- Satellite Name
- Launch Date
- Decay Date (if applicable)
- Classification (if applicable)

Each satellite receives a unique ID, centrally given to all satellites launched. It is a continuous number starting at 1. The satellite with ID 1 is the rocket body of Sputnik 1. Sputnik 1 itself holds ID 2.

The data is obtained by a web crawler that parses the HTML page and finds the information needed. There is also the possibility of getting the data by API, but the API is rate-limited. Parsing HTML is inefficient, but still a quicker approach compared to using the API.

4.1 Latency over whole Data Range

The performance of Starlink latency has changed over time. We looked at RIPE Atlas TLS data that has been collected by built-in measurements (i.e., measurements that are continuously running in each individual probe in a fixed time interval).

s

Figure 4.1 shows the history of median latencies from January 2022 until June 2024 for Germany, the United States, Poland, and Austria⁵.

One can see that the median latency is usually at around 100 to 150 ms for most countries. However, in 2022 the latency was lower compared to 2023. Most of the countries observed show an upwards trend in the late months of 2022 (most of the time in December). In the late months of 2023, the latency starts to decline again. In the last couple of months, we observe an increase in latency once again. We assume the reason for a rise of latency in the congestion of the Starlink network. In the recent time Starlink extended their availability across more countries allowing more users, especially those in countries with less networking infrastructure, to access the network. On the contrary, Starlink also launched more satellites (2022: 3481, 2024: 5700+) and added more PoP which likely reduces the congestion.

Figure 4.7 shows the median and average latencies in European countries. It shows that the north west of Europe experiences the best latencies, likely due to the presence of various GSs. The southern and eastern european countries experience worse latencies. Especially Greece has a high median latency. A cause could be the absence of GSs in the east-european region. Italy on the other hand side experiences a high average latency, even with GSs being present in the country.

Looking at the CDF of Canada results in a similar observation. Figure 4.3 illustrates the CDF plots for 2022 – 2024 for Canada. We observe similar results for other countries, but only chose Canada as it holds sufficient data for a conclusion.

We observe and similar performance of 2022 and 2024, but a drop in performance in 2023, similar to the conclusion we drew before.

Additionally, we conclude that approximately half of the measurement results are below 100 ms, while the other half moves above 100 ms. This opposes research

⁵ These countries were chosen due to the completeness of data

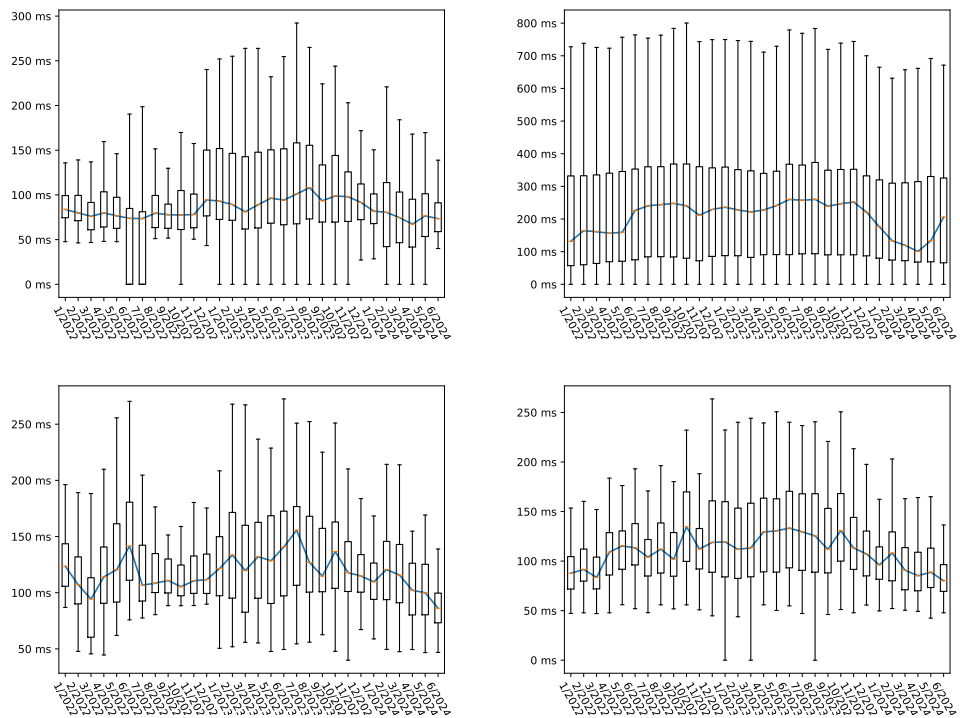


Figure 4.1: Latency History for Germany, the USA, Poland, and Austria in the period 01/2022 to 06/2024

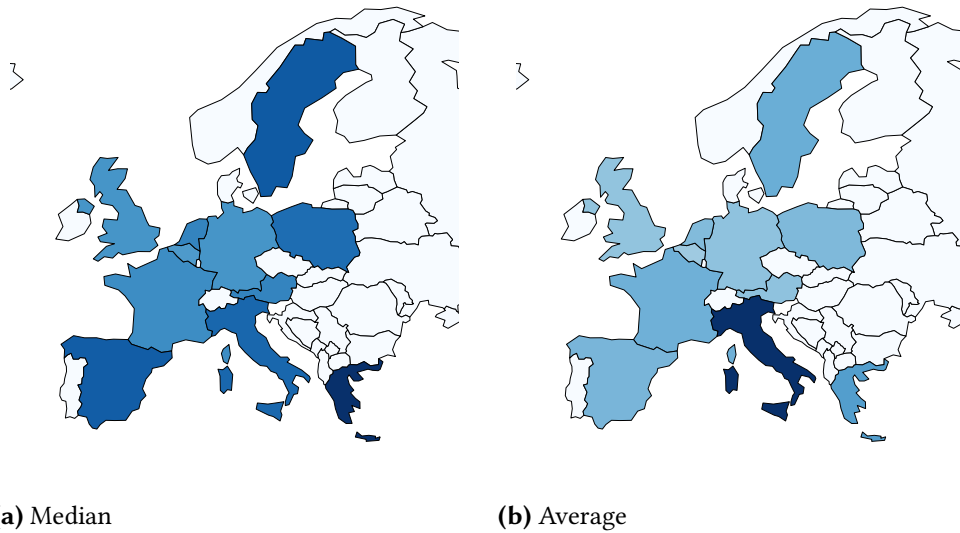


Figure 4.2: Heatmap of Median and Average Latencies in 2024 in Europe.




Figure 4.3: CDF of Starlink Latencies each year in Canada.




Figure 4.4: CDF of Starlink Latencies each year in the USA.

suggesting Starlink performance is mostly in the sub-100 ms area [LLL20; Mic+22; Moh+24; Ram+23].

The CDF has is continuous, up to a specific point, where it flattens, followed by a stronger increase once again. This is similarly observed for curves of other countries (e.g., for the USA in Figure 4.4). This observation suggests that there is a region of latencies Starlink does not serve equally to other latencies. The specific location of the flattening behavior varies between countries, but is usually located between 150 and 250 ms.




Figure 4.5: EquiWidth Histogram of Starlink Latencies 2022 (left), 2023 (middle), and 2024 (right).

Analyzing the flattening behavior in more detail, we used an EquiWidth histogram⁶ to plot the most frequent latencies. Figure 4.5 shows the histogram for the USA. It becomes apparent that there is actually a major gap between the latencies. This gap appears for every country, for every year analyzed. It does not matter how large the country.

As this behavior appears to be independent of time and location, it must originate from Starlink itself. As we lack the possibility to measure other networked satellite ISPs, we do not know whether this is a general property of networked satellite internet.

4.2 Latency Development on Individual Days

We used the data to look at the latency development over a single data for individual days. First, TLS data was used to analyze individual days. However, it appeared that the amount of TLS measurements per hour was insufficient to make a reasonable statement. An example of such a plot for July 9, 2024 is illustrated in Figure 6.1 in the appendix⁷.

Instead, we chose to use Ping data for the analysis of individual days. Ping measurements are not representative for absolute performance metrics [Pel+13], but show a pattern.

4.3 Latency Development on Individual Months

We looked at the latencies within single months.

4.3.1 Latencies on Different Weekdays

It is in question whether latencies vary at different days of the week. For example, the latency might be different comparing Mondays and Sundays. In the case of Starlink, we were not able to find such a pattern. Figure 4.6 illustrates the latencies that occur in each weekday.

One can see that there is no clear pattern that differentiates the individual weekdays from the others. This does not change over the course of the years.

⁶ A histogram that puts all values in a pre-defined number of bins, where all bins cover an equally wide data-range.

⁷ The day was chosen at random. It holds more data than most of the other days.




Figure 4.6: Latencies from 2022 to 2024 per Weekday.

To look in more detail on the data, we looked at the data in specific statistical aggregates (number of measurements, median, average, maximum, and minimum latency). The results are shown in Table 4.1.

They show that the latency does not vary on the individual days. Therefore, you will not have a worse latency on working days (Mon. – Fri.) compared to the weekend (Sat. and Sun.).

However, the table also allows for an interesting comparison between the years 2022, 2023, and 2024. In 2022, the median latency was between 81 and 84 ms. The average latency was about 20 ms higher. Peak latencies achieved up to 43 ms. In 2023, the median latencies were significantly higher at 94 to 98 ms. The average latency was in the worst case 14 ms higher, which indicates a more stable connection, even if the performance was worse compared to 2022. However, 2023 achieved much better peak performance at 26 ms. In 2024, Starlink achieved similar median and average latencies compared to 2022, but also achieving the peak performances from 2023.

Takeaway: Peak and Average Latency since 2022.

Starlink has managed to improve their peak latency compared to 2022 by nearly 20 ms while maintaining their median and average latencies.

Figure 4.7 shows the median and average latencies in European countries. It shows that the north-west of Europe experiences the best latencies, likely due to the presence of various GSs. The southern and eastern European countries experience worse latencies. Especially Greece has a high median latency. A cause could be the absence of GSs in the east European region. Italy on the other hand side experiences a high average latency, even with GSs being present in the country.

Table 4.1: Weekday Statistics in Germany

	Mon.	Tue.	Wed.	Thu.	Fri.	Sat.	Sun.
2022							
#Measurements	899	895	904	903	906	918	893
Median	84	82	82	84	83	81	83
Average	100	100	106	97	99	99	93
Maximum	1211	3090	3056	703	1229	1106	672
Minimum	47	43	46	47	47	47	46
2023							
#Measurements	1448	1447	1443	1429	1430	1436	1462
Median	95	98	94	95	94	95	96
Average	111	111	108	109	107	108	108
Maximum	1245	1227	1233	707	1088	1220	1052
Minimum	28	26	27	27	27	27	27
2024							
#Measurements	914	930	904	913	914	912	918
Median	80	81	80	79	81	80	81
Average	97	93	104	105	95	95	95
Maximum	3147	1592	4374	3624	4368	1230	1320
Minimum	26	26	26	26	26	27	26

4.4 Disconnect Events of RIPE Atlas Probes

Probes in RIPE Atlas automatically monitor the events when they are disconnected from RIPE Atlas and when they re-connect again. Those disconnect events give evidence about the state of a network. Frequently occurring disconnect events might indicate an unhealthy network connection.

We analyzed the number of disconnect events for the Starlink probes in RIPE Atlas. Figure 4.8 shows the number of disconnect events per month in the time from January 2022 to July 2024.

One sees variously different patterns. While in Germany and the United Kingdom very little disconnect events appear for the most time, the United States show a relatively high number of disconnect events, increasing over time. The USA experiences higher levels of disconnect events in June 2024, January 2023, and May to August 2022. A clear pattern is not visible here. On the other side, Germany experiences a high number of disconnect events in May to August 2022 (which also

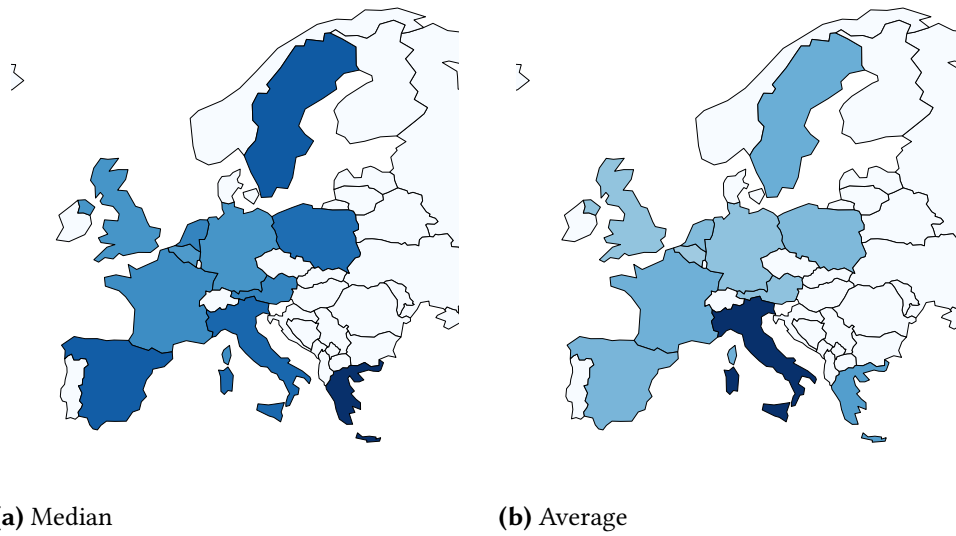


Figure 4.7: Heatmap of Median and Average Latencies in 2024 in Europe (darker color means a higher latency).

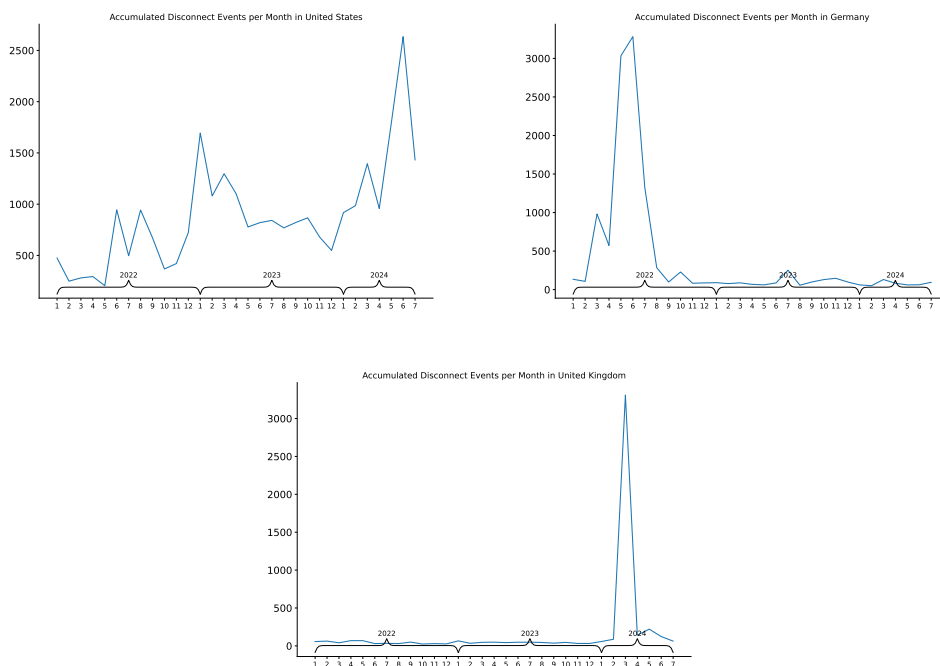


Figure 4.8: Number of Disconnect Events in the USA, Germany, and the United Kingdom from January 2022 to July 2024.




Figure 4.9: Accumulated Number of Disconnect Events per Month in all Countries

happens in the USA), while the United Kingdom does not show such a behavior. The UK experiences a higher number of disconnect events in March 2024, but never before that. There is no clear pattern observable, when comparing those three countries. A similar behavior was observed for other countries.

Still, accumulating the number of disconnect events over all countries show significant spikes. Figure 4.9 illustrates the pattern. Especially September and October 2023, the beginning of 2024 and June 2024 show a much higher frequency. May and June 2022 have a significant spike.

We have to note that the data shown in Figure 4.9 does not allow making conclusions. It holds different numbers of probes over time, so an increase in disconnect events is expected. To draw conclusions for Starlink behavior, the data needs to be normalized. Chapter 4.4.1 explains this problem in more detail.

Takeaway: Usage of Disconnect Events

Accumulating Disconnect Events of Starlink probes seem to show a behavior that causes time intervals with higher occurrences. However, normalization is required to verify the observation. Normalization would be possible by using the active number of probes per time interval, but this was not done as it is not an easy task.

4.4.1 Correlation of Disconnect Events with other Metrics

One could correlate the occurrences of disconnect events with other metrics (e.g., latencies and packet loss). Usually, this is done by comparing two metrics with each other, where each metric is normalized. By normalizing, two values of the same metric are comparable. In the case of disconnect events, the number of disconnect events per time interval grows with the number of probes connected. Therefore, the number of occurrences would be normalized by number of probes connected.

However, the number of probes connected within a time interval is hard to find. Some probes might be connected for a longer time than others dominating the occurrences of disconnect events. Others might be connected for a very short interval with many disconnects. Both scenarios compromise the normalization of the disconnect event occurrences, leading to not meaningful results.⁸

Therefore, we did not proceed further in correlating disconnect events with other metrics.

4.5 Latency and Packet Loss Correlation

Latency and packet loss are closely bonded. An increase of packet loss might yield an increase in latency and vice-versa. For the data presented in the previous chapter, we looked at the interval from January 2022 to June 2024. For each month, we used the overall packet loss and the median latency. We used a Pearson, Kendall, and Spearman correlation to determine possible correlations between latency and packet loss. The results are shown in Table 4.2.

In a correlation, a value close to 1 or -1 shows a strong correlation. Values close to 0 do not appear to be correlated. In the results, some countries show a correlation, e.g., France or Czechia. Other countries, e.g., Canada or Australia do not show correlating values. As the number of countries with a stronger and those with a weaker correlation is equally distributed, we cannot conclude a correlation between latency and packet loss in Starlink networks.

4.5.1 Correlation in Single Years

Analyzing single years give a better understanding of the development of correlation in the year 2022, 2023, and 2024 (until June). The Tables 4.3, 4.4, and 4.5 show the

⁸ It might be well-possible to find a normalization method for disconnect events, but we did not invest in this direction.

Table 4.2: Packet Loss and Latency Correlation

Country January 2022 - June 2024	Pearson Correlation	Kendall Correlation	Spearman Correlation
Falkland Islands (Malvinas)	0.447641	0.595238	0.624282
Réunion	0.636822	0.716314	0.847658
Kiribati	0.819596	0.794851	0.801966
Canada	-0.090981	0.218391	0.280534
Poland	0.093155	-0.025287	-0.073637
Haiti	0.383601	0.670768	0.844964
Spain	-0.323940	-0.163650	-0.179434
Czechia	0.700248	0.730392	0.806351
United States	-0.556997	-0.425287	-0.610234
France	0.893576	0.664368	0.840267
Italy	-0.073752	0.181609	0.293882
United Kingdom	-0.238013	-0.227586	-0.279644
Honduras	0.415387	0.410256	0.575585
Australia	-0.025590	0.013841	0.035834
Netherlands	0.190923	0.055364	0.111952
Greece	0.249727	0.496864	0.668590
Sweden	0.678500	0.726984	0.896199
Austria	0.054429	0.039080	0.073192
Belgium	-0.082786	-0.029885	-0.031368
Philippines	0.254523	0.460118	0.676049
Virgin Islands, U.S.	0.712898	0.706384	0.878571
Germany	-0.787348	-0.549425	-0.753059

individual correlation values for the year. There are only countries displayed, where a correlation value could be calculated.

In 2022, there was little correlation at all. Likely, other factors like satellite infrastructure or presence of PoP might have been more contributing to performance issues. 2023 showed a stronger correlation compared to 2022. We assume that this is due to a more stable network. The interaction between packet loss and latency might have become more significant. In 2024, the correlation decreased once again slightly. This behavior is similar to the one observed for the latencies in recent months. A possible explanation for that is a growing user base of Starlink leading to a congestion of the system. It has to be noted that the growth of number of satellites did not increase from 2023 to 2024 as much as it did from 2022 to 2023, which might be a cause of the issue.

Table 4.3: Packet Loss and Latency Correlation in 2022

Country 2022	Pearson Correlation	Kendall Correlation	Spearman Correlation
Canada	0.478988	0.272727	0.307692
Poland	0.601734	0.303030	0.342657
Spain	-0.286931	-0.053571	-0.037594
United States	-0.241717	-0.212121	-0.363636
France	0.177126	-0.060606	-0.083916
Italy	0.646487	0.484848	0.587413
United Kingdom	0.514397	0.363636	0.496503
Honduras	0.208592	-0.333333	-0.426573
Australia	0.870466	0.325669	0.450715
Netherlands	0.159295	-0.106873	-0.164624
Greece	0.811845	0.836660	0.853766
Austria	0.013609	-0.060606	-0.146853
Belgium	-0.283806	-0.242424	-0.272727
Germany	-0.663075	-0.303030	-0.517483

Table 4.4: Packet Loss and Latency Correlation in 2023

Country 2023	Pearson Correlation	Kendall Correlation	Spearman Correlation
Falkland Islands (Malvinas)	0.364389	0.466667	0.518072
Réunion	0.587427	0.467801	0.592125
Canada	0.370590	0.363636	0.608392
Poland	-0.261476	-0.090909	-0.216783
Haiti	0.757392	0.390673	0.529108
Spain	-0.413832	-0.160714	-0.218045
Czechia	0.415844	0.383333	0.458333
United States	-0.302660	-0.242424	-0.251748
France	0.779247	0.666667	0.818182
Italy	-0.568765	-0.272727	-0.412587
United Kingdom	-0.812772	-0.606061	-0.748252
Honduras	0.343476	0.516667	0.617754
Australia	-0.578546	-0.121212	-0.258741
Netherlands	0.212359	0.181818	0.314685
Greece	-0.268072	0.212121	0.195804
Sweden	0.515385	0.533333	0.717391
Austria	-0.294406	-0.121212	-0.125874
Belgium	0.441123	0.454545	0.573427
Philippines	-0.196523	-0.030303	-0.034965
Virgin Islands, U.S.	0.553108	0.578196	0.783080
Germany	-0.620018	-0.424242	-0.622378

Table 4.5: Packet Loss and Latency Correlation in 2024

Country 2024	Pearson Correlation	Kendall Correlation	Spearman Correlation
Réunion	0.301680	-0.066667	-0.085714
Kiribati	0.744916	0.673575	0.718421
Canada	-0.003851	-0.600000	-0.771429
Poland	0.239930	0.333333	0.371429
Haiti	0.376160	0.200000	0.257143
Spain	-0.044513	-0.333333	-0.142857
United States	-0.529555	-0.600000	-0.771429
France	0.418669	0.200000	0.257143
Italy	-0.600898	-0.866667	-0.942857
United Kingdom	0.243811	0.333333	0.428571
Honduras	0.431954	0.333333	0.542857
Australia	-0.155578	-0.066667	-0.085714
Netherlands	0.433624	0.466667	0.542857
Greece	-0.813812	-0.466667	-0.600000
Sweden	0.280194	0.066667	0.085714
Austria	0.915449	0.466667	0.600000
Belgium	0.428515	0.333333	0.371429
Philippines	-0.920833	-0.466667	-0.428571
Virgin Islands, U.S.	-0.450938	-0.333333	-0.428571
Germany	-0.603772	-0.333333	-0.485714

4.5.2 Correlation with the Number of Probes

It is possible that the data is insufficient resulting in a correlation between latency and packet loss being invisible. A possibility is to look at the number of probes being available for each country. Therefore, we took the data from Table 4.2 and correlated it with the number of probes available for RIPE Atlas in each country. The resulting table is shown in Figure 4.6.

We correlated the correlation values with the number of probes per country. This resulted in the following correlation values:

- Pearson Correlation: ≈ -0.44
- Kendall Correlation: ≈ -0.44
- Spearman Correlation: ≈ -0.53

As the correlation does not reach values close to 0, 1, or -1, we cannot conclude a correlation with the number of probes.

Table 4.6: Packet Loss, Latency and Number of Probes Correlation

Country Jan. 2022 - June 2024	Pearson Correlation	Kendall Correlation	Spearman Correlation	Number of Probes
Falkland Islands (Malvinas)	0.447641	0.595238	0.624282	1
Réunion	0.636822	0.716314	0.847658	1
Kiribati	0.819596	0.794851	0.801966	2
Canada	-0.090981	0.218391	0.280534	11
Poland	0.093155	-0.025287	-0.073637	1
Haiti	0.383601	0.670768	0.844964	3
Spain	-0.323940	-0.163650	-0.179434	4
Czechia	0.700248	0.730392	0.806351	1
United States	-0.556997	-0.425287	-0.610234	53
France	0.893576	0.664368	0.840267	18
Italy	-0.073752	0.181609	0.293882	4
United Kingdom	-0.238013	-0.227586	-0.279644	11
Honduras	0.415387	0.410256	0.575585	1
Australia	-0.025590	0.013841	0.035834	8
Netherlands	0.190923	0.055364	0.111952	2
Greece	0.249727	0.496864	0.668590	1
Sweden	0.678500	0.726984	0.896199	1
Austria	0.054429	0.039080	0.073192	4
Belgium	-0.082786	-0.029885	-0.031368	2
Switzerland	NaN	NaN	NaN	1
Philippines	0.254523	0.460118	0.676049	3
Benin	NaN	NaN	NaN	2
Virgin Islands, U.S.	0.712898	0.706384	0.878571	1
Germany	-0.787348	-0.549425	-0.753059	10

Takeaway: The correlation of latency and packet loss.

Starlink does not show a clear pattern that indicates a correlation between latency and packet loss. However, this is highly dependent on the country of origin.

4.6 Traceroute Analysis

Looking at traceroute results, we can make conclusions about the routing behavior of Starlink network devices. In the data we gathered more than forty million traceroute measurements. Those include primarily built-in measurements from




Figure 4.10: Visualization of a Traceroute from Kiribati

RIPE Atlas probes to **.root-servers.net* for Starlink probes (Autonomous System Number (ASN) 14593).

Routing Behavior

First, we found that the satellite hops are likely invisible to the traceroute. We conclude that as there are no hops visible above water, even for probes located in remote regions, e.g., Kiribati in the Pacific Ocean. In Figure 4.10, one can see a visualization of a traceroute result from Kiribati to *f.root-servers.net*. One can see that the first visible hop is located in New Zealand or an island to the north of New Zealand. However, if satellites were visible, we'd be able to observe more satellites.

Coming from the first insight, we can also conclude that Inter-Satellite Links (ISLs) are enabled. If they were not, we would likely not be able to see a successful traceroute from Kiribati to a location. The next closest known PoP is on Hawaii. However, the distance between both is 4000 km, which is more than a single satellite can cover. Aside from that, we do not observe the usage of the PoP in Hawaii, but in more distant locations, which just strengthens the argument. Therefore, we conclude that ISLs are enabled. This is special interest, as it was not clear in recent research [Hau+20].

Verify that this poses actually a privacy concern / Remove if there is no proof or no privacy concern

4.6.1 Privacy Concerns in Traceroute Data

One of the most important responsibilities of an Internet Service Provider (ISP) is to ensure the privacy of its users. This also includes to route traffic only in trusted countries. In the case of Starlink, we were able to observe a different behavior. We

Table 4.7: IP Hitlist for Built-In Traceroute Measurements

Hits	City	Country	Organization	IP Address
10634	Frankfurt am Main	Germany	AS1299	62.115.37.20
8202	Offenbach	Germany	Unknown	80.81.192.154
6207	Amsterdam	Netherlands	Unknown	193.239.116.217
5582	Frankfurt am Main	Germany	AS2914	213.198.72.18
5257	Frankfurt am Main	Germany	AS3257	89.149.137.14
4932	Chicago	United States	AS14593	206.224.65.178
4916	Chicago	United States	AS14593	206.224.65.180
4850	Chicago	United States	AS14593	206.224.65.182
4755	Chicago	United States	AS14593	206.224.65.184
4358	Miami	United States	AS49791	81.31.213.126
4333	Zürich	Switzerland	Unknown	185.1.147.30
4256	Tokyo	Japan	Unknown	210.173.176.242
4179	Chicago	United States	AS14593	206.224.65.186
4035	Chicago	United States	AS14593	206.224.65.192
4014	Frankfurt am Main	Germany	AS6762	213.144.184.30
4010	Chicago	United States	AS14593	206.224.65.190
4005	Chicago	United States	AS14593	206.224.65.188
4002	Frankfurt am Main	Germany	AS1299	62.115.124.118
3990	Singapore	Singapore	AS2497	202.232.1.69
3827	Frankfurt am Main	Germany	AS6939	72.52.92.70

looked at a slice of the built-in traceroute measurements from German Starlink probes and analyzed their most common targets. We filtered for anycasted servers (e.g., *.root-servers.net) and bogon IPs (i.e., IPs that cannot be associated with metadata).

In Table 4.7, the top twenty most frequent hits of IP addresses are shown. The IP addresses are joined with data from IPinfo. As traffic goes from a German probe to an anycasted server, located in Germany, one would expect little traffic outside Germany, and none outside Europe. The top five IP addresses are located within or close to Germany, but the next five already involve traffic to the United States. Here, we observe an unexpected behavior. Assuming that the data is not flawed, this is a clear violation of guiding privacy principles.

5

Conclusions & Outlook




Figure 6.1: Latency over the July 9, 2024 using TLS data.

List of Acronyms

ISLs Inter-Satellite Links	28
ISL Inter-Satellite Link	4
ISP Internet Service Provider	28
LEO Low-Earth Orbit	4
MEO Medium-Earth Orbit	4
GEO Geostationary Orbit	4
ASN Autonomous System Number	28
PoP Point of Presence	3
GS Ground Station	3
CFR Cloudflare Radar	12
IQI Internet Quality Index	12

Bibliography

- [Clo] Cloudflare. *Cloudflare Radar Docs Glossary: Internet Quality Index* (see page 12).
- [Clo24] Cloudflare. *The Global Cloudflare-Network*. Cloudflare Website. Sept. 2024 (see page 13).
- [Dav23] Carlos Rodrigues David Belson. **Introducing the Cloudflare Radar Internet Quality Page**. *The Cloudflare Blog* (June 2023). URL: <https://blog.cloudflare.com/introducing-radar-internet-quality-page/> (see page 12).
- [Hau+20] Yannick Hauri, Debopam Bhattacharjee, Manuel Grossmann, and Ankit Singla. **"Internet from Space" without Inter-satellite Links**. In: *HotNets '20: The 19th ACM Workshop on Hot Topics in Networks, Virtual Event, USA, November 4-6, 2020*. Ed. by Ben Y. Zhao, Heather Zheng, Harsha V. Madhyastha, and Venkat N. Padmanabhan. ACM, 2020, 205–211. DOI: 10.1145/3422604.3425938. URL: <https://doi.org/10.1145/3422604.3425938> (see page 28).
- [Kas+20] Simon Kassing, Debopam Bhattacharjee, André Baptista Águas, Jens Eirik Saethre, and Ankit Singla. **Exploring the "Internet from space" with Hypatia**. In: *IMC '20: ACM Internet Measurement Conference, Virtual Event, USA, October 27-29, 2020*. ACM, 2020, 214–229. DOI: 10.1145/3419394.3423635. URL: <https://doi.org/10.1145/3419394.3423635> (see page 6).
- [Koh24] Thomas Kohnmann. **Everything you need to know about Project Kuiper, Amazon's satellite broadband network**. June 2024 (see page 6).
- [LLL20] Zeqi Lai, Hewu Li, and Jihao Li. **StarPerf: Characterizing Network Performance for Emerging Mega-Constellations**. In: *28th IEEE International Conference on Network Protocols, ICNP 2020, Madrid, Spain, October 13-16, 2020*. IEEE, 2020, 1–11. DOI: 10.1109/ICNP49622.2020.9259357. URL: <https://doi.org/10.1109/ICNP49622.2020.9259357> (see pages 6, 17).
- [Ma+23] Sami Ma, Yi Ching Chou, Haoyuan Zhao, Long Chen, Xiaoqiang Ma, and Jiangchuan Liu. **Network Characteristics of LEO Satellite Constellations: A Starlink-Based Measurement from End Users**. In: *IEEE INFOCOM 2023 - IEEE Conference on Computer Communications, New York City, NY, USA, May 17-20, 2023*. Movement, Environment, LEO, Performance, Comparison LEO vs Terrestrial , TCP vs UDP, Power Consumption, Bent-Pipe. IEEE, 2023, 1–10. DOI: 10.1109/INFOCOM53939.2023.10228912. URL: <https://doi.org/10.1109/INFOCOM53939.2023.10228912> (see page 4).

- [Mic+22] François Michel, Martino Trevisan, Danilo Giordano, and Olivier Bonaventure. **A first look at starlink performance**. In: *Proceedings of the 22nd ACM Internet Measurement Conference, IMC 2022, Nice, France, October 25-27, 2022*. Ed. by Chadi Barakat, Cristel Pelsser, Theophilus A. Benson, and David R. Choffnes. ACM, 2022, 130–136. DOI: [10.1145/3517745.3561416](https://doi.org/10.1145/3517745.3561416). URL: <https://doi.org/10.1145/3517745.3561416> (see page 17).
- [Moh+24] Nitinder Mohan, Andrew E. Ferguson, Hendrik Cech, Rohan Bose, Prakita Rayyan Renatin, Mahesh K. Marina, and Jörg Ott. **A Multifaceted Look at Starlink Performance**. In: *Proceedings of the ACM on Web Conference 2024, WWW 2024, Singapore, May 13-17, 2024*. Ed. by Tat-Seng Chua, Chong-Wah Ngo, Ravi Kumar, Hady W. Lauw, and Roy Ka-Wei Lee. ACM, 2024, 2723–2734. DOI: [10.1145/3589334.3645328](https://doi.org/10.1145/3589334.3645328). URL: <https://doi.org/10.1145/3589334.3645328> (see pages 3, 4, 17).
- [PA19] Timothy Pratt and Jeremy E Allnutt. **Satellite communications**. John Wiley & Sons, 2019 (see page 3).
- [Pel+13] Cristel Pelsser, Luca Cittadini, Stefano Vissicchio, and Randy Bush. **From Paris to Tokyo: on the suitability of ping to measure latency**. In: *Proceedings of the 2013 Internet Measurement Conference, IMC 2013, Barcelona, Spain, October 23-25, 2013*. Ed. by Konstantina Papagiannaki, P. Krishna Gummadi, and Craig Partridge. ACM, 2013, 427–432. DOI: [10.1145/2504730.2504765](https://doi.org/10.1145/2504730.2504765). URL: <https://doi.org/10.1145/2504730.2504765> (see page 18).
- [Ram+23] Aravindh Raman, Matteo Varvello, Hyunseok Chang, Nishanth Sastry, and Yasir Zaki. **Dissecting the Performance of Satellite Network Operators**. *PACMNET 1:CoNEXT3* (2023), 15:1–15:25. DOI: [10.1145/3629137](https://doi.org/10.1145/3629137). URL: <https://doi.org/10.1145/3629137> (see pages 4, 17).
- [roo24] root-servers.org. *Root Servers*. Website. Sept. 2024. URL: <https://root-servers.org/> (see page 13).
- [Seg20] Sascha Segan. **Tested: SpaceX’s Starlink Satellite Internet Service Is Fast, But It’ll Cost You**. *PC Mag* (Oct. 2020) (see page 4).
- [Ste+24] Aleksandr Stevens, Blaise Iradukunda, Brad Bailey, and Ramakrishnan Durairajan. **Can LEO Satellites Enhance the Resilience of Internet to Multi-hazard Risks?** In: *Passive and Active Measurement - 25th International Conference, PAM 2024, Virtual Event, March 11-13, 2024, Proceedings, Part II*. Ed. by Philipp Richter, Vaibhav Bajpai, and Esteban Carisimo. Vol. 14538. Lecture Notes in Computer Science. Springer, 2024, 170–195. DOI: [10.1007/978-3-031-56252-5_9](https://doi.org/10.1007/978-3-031-56252-5_9). URL: https://doi.org/10.1007/978-3-031-56252-5_9 (see page 5).

Declaration of Authorship

I hereby declare that this thesis is my own unaided work. All direct or indirect sources used are acknowledged as references.

Potsdam, September 23, 2024

Robert Alexander Uwe Richter